

MANAGEMENT AND PROCESSING OF BIG DATA

LEVEL-I

SESSION-2



Recap

- Data processing trend
- Big Data and its characteristics
- Applications of big data
- Various computing technologies
- History of Hadoop
- Major components of Hadoop cluster
- Unix commands
- HDFS commands

Agenda for today

- The Hadoop Distributed File System
- MapReduce detailed discussion
- Running your first MapReduce program
- Automation of big data jobs

HDFS

- HDFS is a file system designed for storing very large files with streaming data access patterns, running on clusters of commodity hardware

Large Files

Streaming data
access

Commodity
hardware

HDFS Blocks

- Single unit of storage
- Default block size is 128MB
- Size of block will drive the ratio of time to read a block to the seek for a block

Exercise


- What should be the block size to make seek time 1% of read time for given hardware configuration
 - Seek time: 10 ms
 - Data read rate: 100MB/s
- Solution:
 - Let say x MB is the block size then read time = $x/100$ seconds
 - To fulfill the given condition 1% of $x/100 = 10$ ms
 - Hence $x=100$ MB

Benefits of blocks

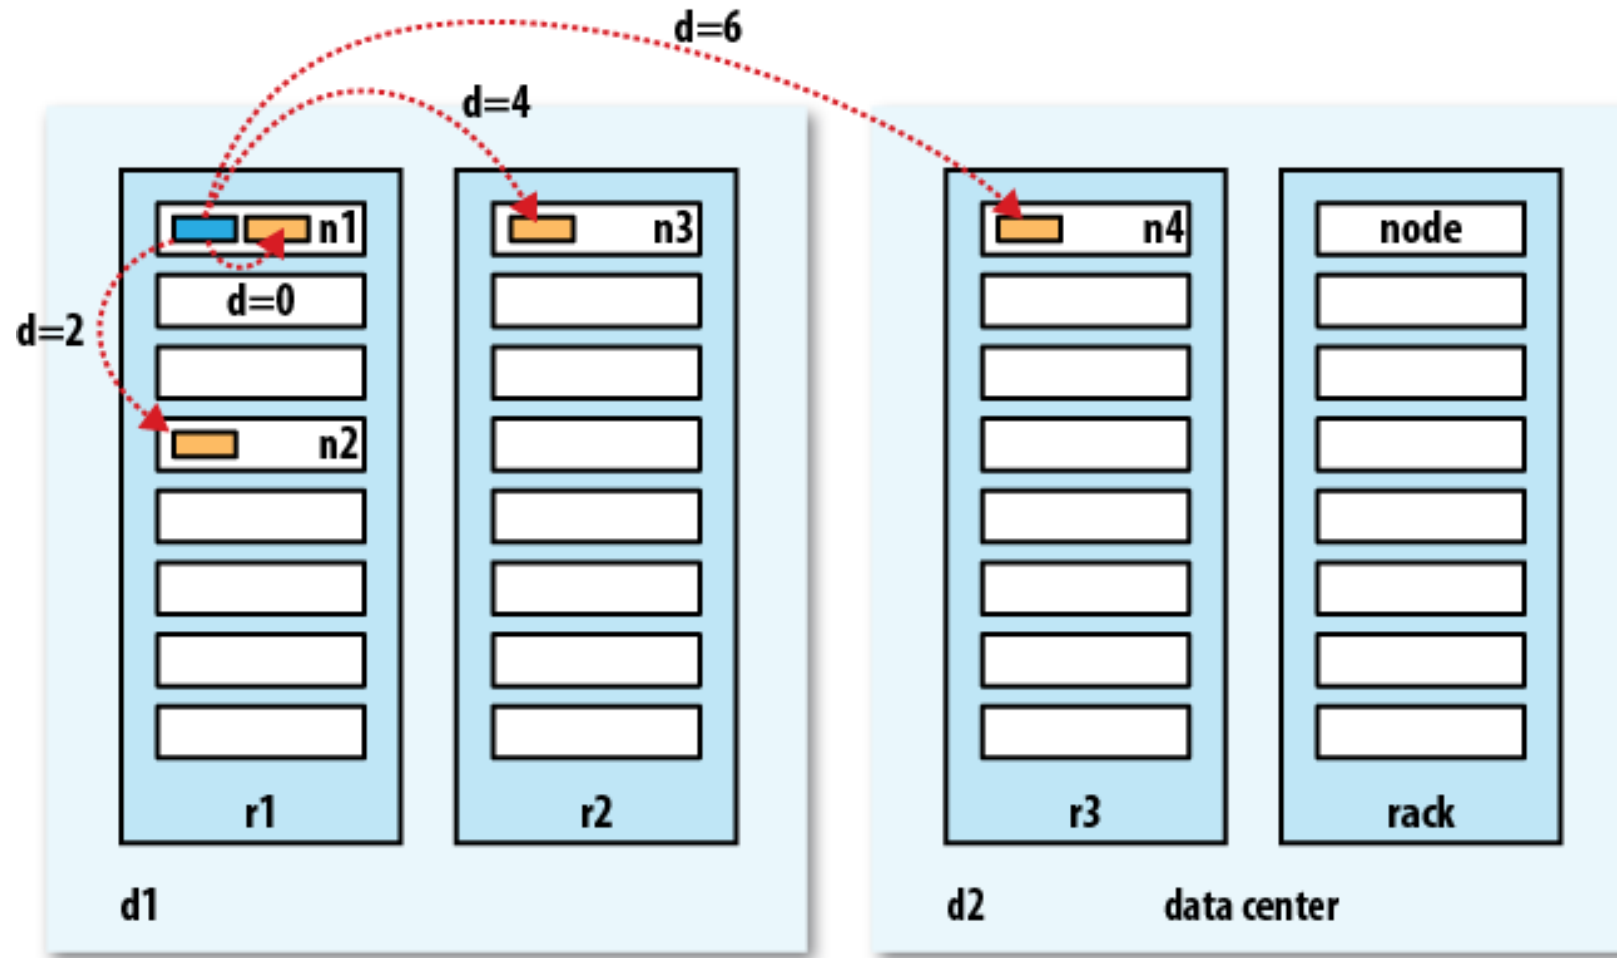
- Files can be larger than a single disk
- Simplicity at storage level as data node doesn't store any metadata
- Fault tolerance by replicating blocks

Filesystem metadata


- Namenode stores the metadata
- Backup of metadata on secondary namenode



Network Topology



Rack awareness



HDFS CLI Read Commands

- Copy a file from local file system to HDFS

`hadoop fs -copyFromLocal <source> <target>`

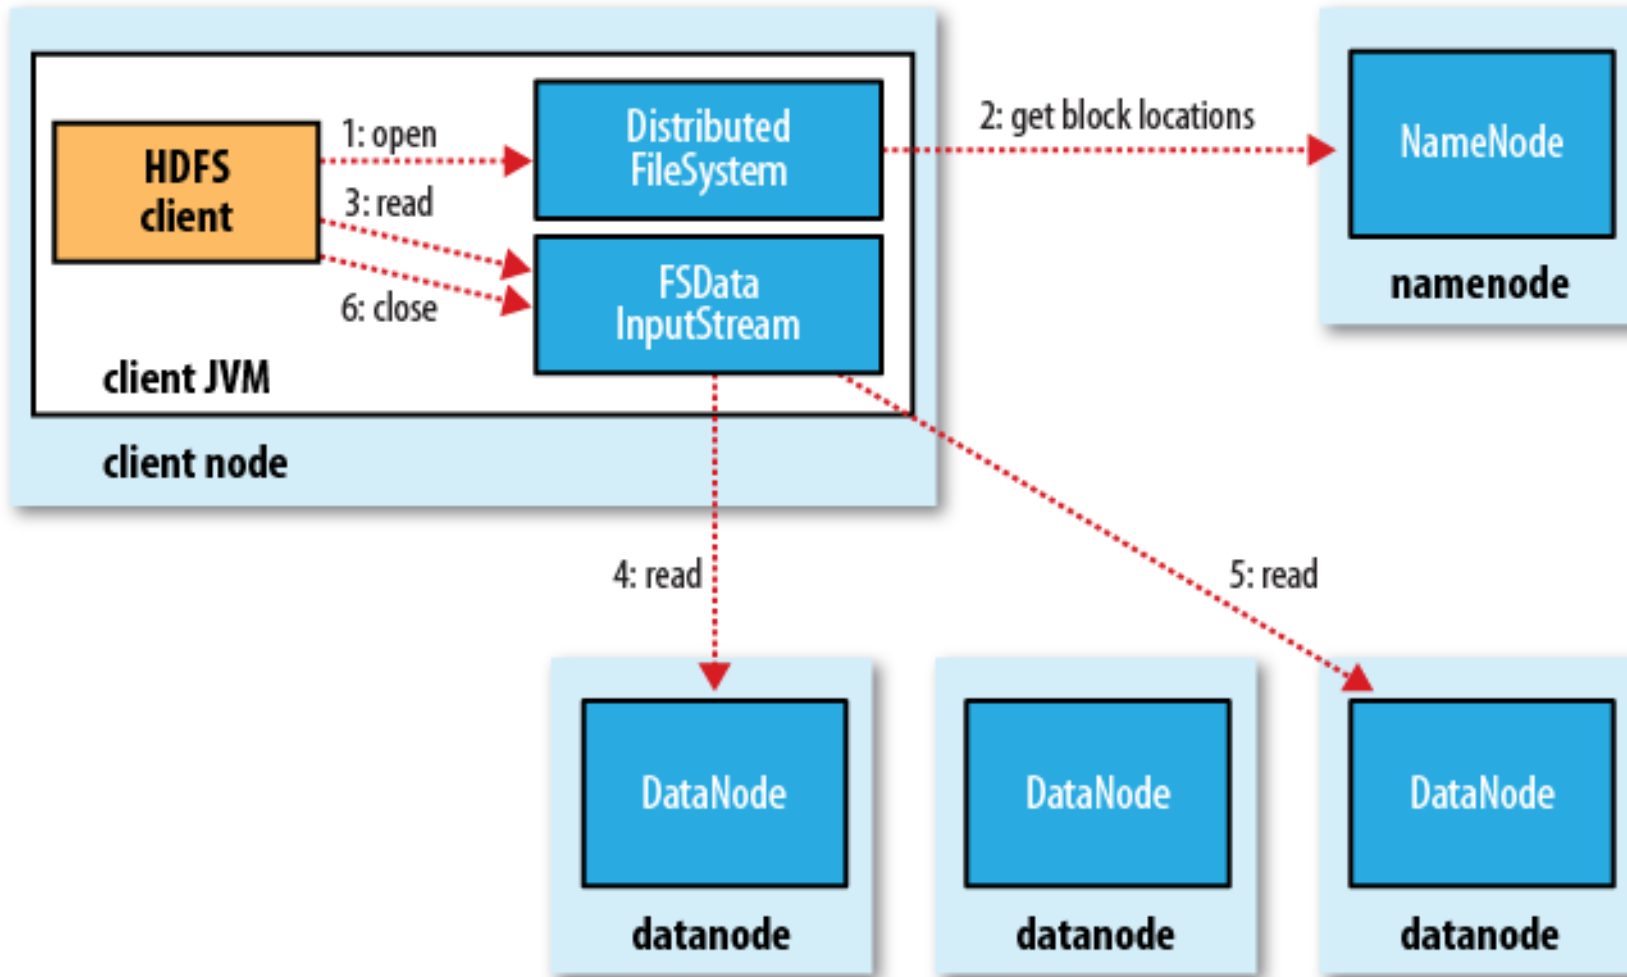
- Get merged file from HDFS to local file system

`hadoop fs -getmerge <source> <target>`

- Cat a file from HDFS

`hadoop fs -cat <filename>`

HDFS Read operation



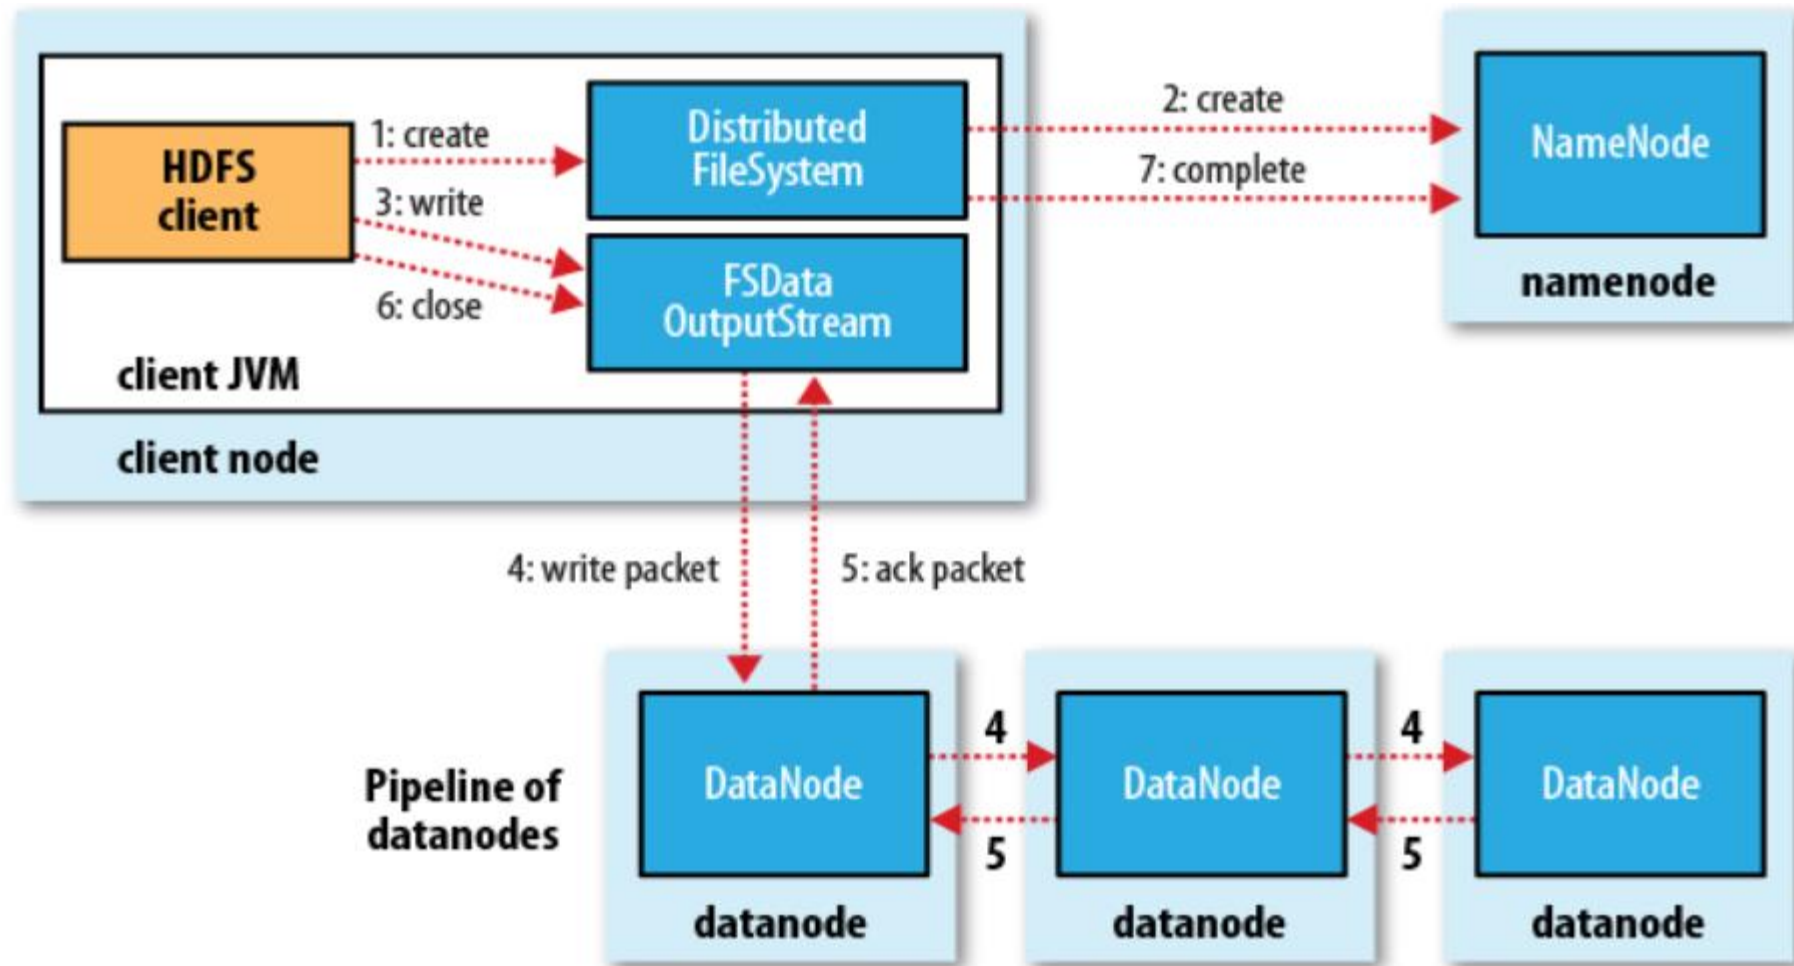
HDFS CLI Write Commands

- Write a file from local file system to HDFS

```
hadoop fs -copyFromLocal <source> <target>
```

```
hadoop fs -put <source> <target>
```

HDFS Write operation



HDFS not made for

- Low-latency data access
- Lots of small files
 - Each file | block | directory stores around 150 bytes of metadata. Hence 1 million files each of one block will consume 300 MB of storage on Namenode
- Multiple writers, arbitrary file modifications

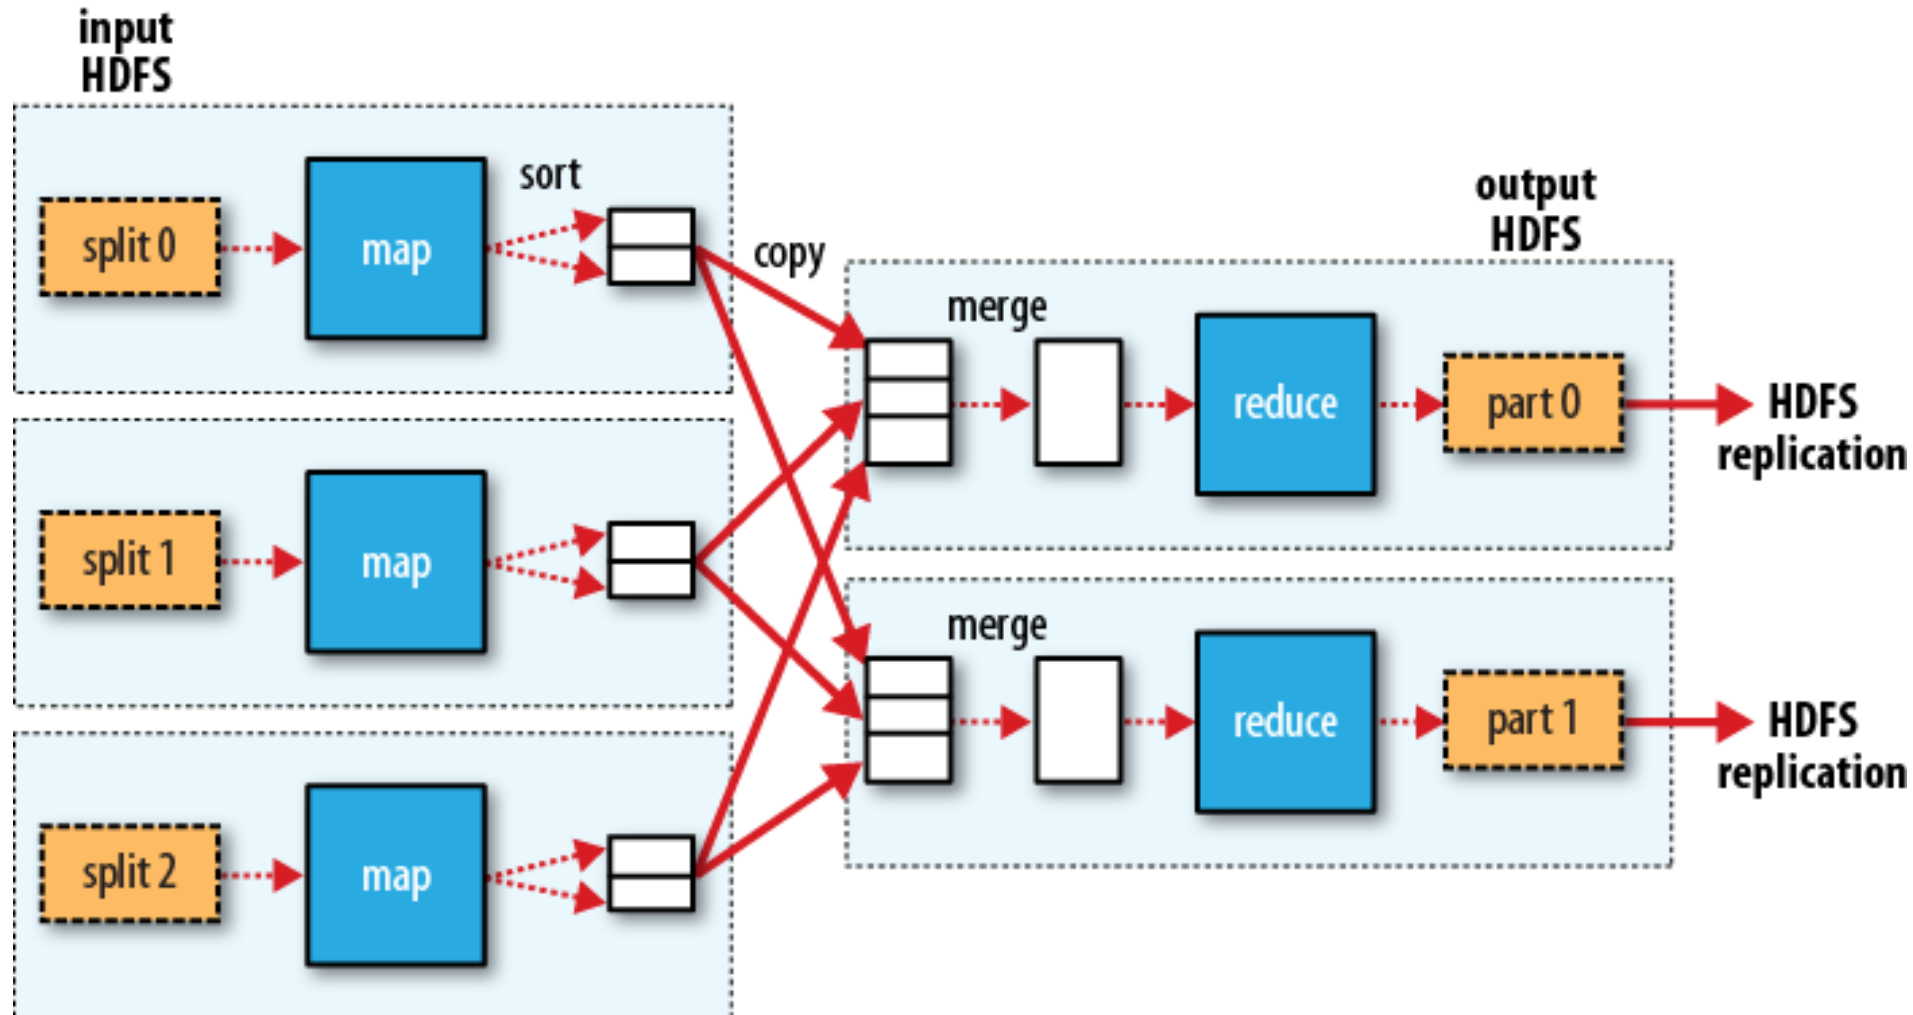
Exercise

- Calculate the memory(RAM) requirement on Namenode for given cluster configurations
 - Cluster size: 200 nodes
 - Storage capacity of each node: 24 TB
 - Block size: 128MB
 - Replication factor: 3
 - Metadata storage size for each block: 150 bytes
- Solution
 - $(200 * 24 * 10^{12} * 150) / (128 * 10^6 * 3)$

MapReduce

- Two major phases: Map and Reduce
- Notion of <Key, Value> pairs
- Divides job into multiple tasks
- Map: extract important information from each record
- Reduce: Aggregate, Summarize, Filter, Transform

MapReduce Stages




Map Tasks

- What is a good approach to decide how many map tasks a job should launch?
- Less number of big tasks
vs
higher number of small tasks
- Normally same as input data blocks


Input Splits

- Blocks are of fixed size
- Good chances of records being split between two block



Task to node mapping

- Notion of data locality



Reduce Tasks

- Can be configured by programmer
- Normally same as #datanodes participating in execution
- Input Key and Value type should be same as output type of mapper
- One output file per reducer under output directory
- Generates exception if output directory already exists.

MapReduce: Mapper code

```
public class WebHitCounterMapper extends
Mapper<Input Key, Input Value, Output Key, Output Value>
{

    public void map(Input Key, Input Value, Context
context) throws IOException, InterruptedException {

        <MAP Logic goes here>

        context.write(Output Key, Output Value)
    }
}
```

MapReduce: Reducer code

```
public class WebHitCounterReducer extends
Reducer<Input Key, Input Value, Output Key, Output Value>
{
    public void reduce(Input Key, Iterable<Value Data type>
values, Context context) throws IOException, InterruptedException {

        <REDUCE logic goes here>

        context.write(Output Key, Output Value);
    }
}
```

MapReduce: Driver Code

```
public class WebHitCounterMain {  
    public static void main(String[] args) throws Exception {  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");  
  
    }  
}
```

MapReduce: Driver Code

```
public class WebHitCounterMain {  
    public static void main(String[] args) throws Exception {  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");  
  
        job.setJarByClass(main.WebHitCounterMain.class);  
        job.setMapperClass(mapper.WebHitCounterMapper.class);  
        job.setReducerClass(reducer.WebHitCounterReducer.class);  
  
    }  
}
```


MapReduce: Driver Code

```
public class WebHitCounterMain {  
    public static void main(String[] args) throws Exception {  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");  
  
        job.setJarByClass(main.WebHitCounterMain.class);  
        job.setMapperClass(mapper.WebHitCounterMapper.class);  
        job.setReducerClass(reducer.WebHitCounterReducer.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
    }  
}
```

MapReduce: Driver Code

```
public class WebHitCounterMain {  
    public static void main(String[] args) throws Exception {  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");  
  
        job.setJarByClass(main.WebHitCounterMain.class);  
        job.setMapperClass(mapper.WebHitCounterMapper.class);  
        job.setReducerClass(reducer.WebHitCounterReducer.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
    }  
}
```

MapReduce: Driver Code

```
public class WebHitCounterMain {  
    public static void main(String[] args) throws Exception {  
  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Daily Web Hit Counter");  
  
        job.setJarByClass(main.WebHitCounterMain.class);  
        job.setMapperClass(mapper.WebHitCounterMapper.class);  
        job.setReducerClass(reducer.WebHitCounterReducer.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
    }  
}
```

Additional Conf Properties

- Configure maximum retries for failed tasks
 - `setMaxMapAttempts(int n)`
 - `setMaxReduceAttempts(int n)`
- Configure reduce tasks explicitly
 - `setNumReduceTasks(int tasks)`

Input formats

Input format	Description
TextInputFormat	Read Text file line by line. Key is offset and value is record text
KeyValueTextInputFormat	Tab separated key values from a text file
SequenceFileInputFormat<K,V>	Hadoop's file format

Output Formats

Output Format	Description
TextOutputFormat<K,V>	Tab separated key value pairs in plain text format. One record per key value pair
SequenceFileOutputFormat<K,V>	Hadoop's Sequence file format
NullOutputFormat<K,V>	Nothing. Helps in map only job

Programming Exercise

Challenges with Hadoop 1

- Applications were limited to MapReduce implementations only
- Namenode machine crash or maintenance activity
- Namespace scaling
- Backup and Recovery
- Batch oriented architecture
- Support for various file formats
- Dual responsibilities of Job tracker

Hadoop 2

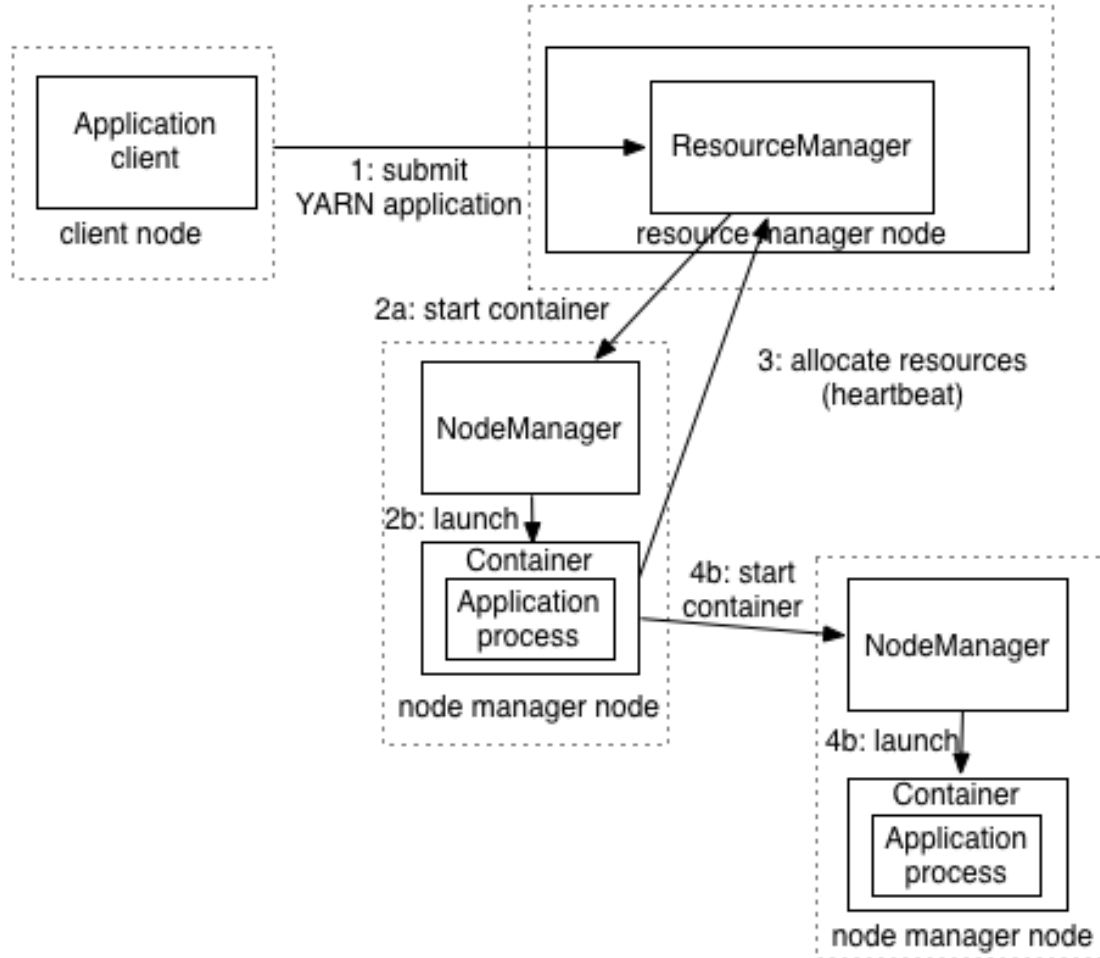
- Support for other data processing engines
- High Availability
- HDFS Federation
- HDFS Snapshot
- Introduced Streaming and Interactive analysis
- Support for various file formats
- YARN

YARN

- Yet Another Resource Negotiator

MapReduce 1	YARN
Job Tracker	Resource Manager, Application Master and Timeline server
Task Tracker	Node Manager
Slot	Containers

YARN model



Pros of YARN

- Scalability
- Availability
- Utilization
- Multitenancy

References

- Reference Book
 - Hadoop: The definitive guide by Tom White ([Weblink](#))
- MapReduce google paper
 - <https://ai.google/research/pubs/pub62>
- Filesystem metadata on Namenode
 - <https://hortonworks.com/blog/hdfs-metadata-directories-explained/>
- Rack awareness
 - <https://data-flair.training/blogs/rack-awareness-hadoop-hdfs/>
 - <https://community.hortonworks.com/articles/43057/rack-awareness-1.html>
- Concurrency vs Parallelism
 - <https://stackoverflow.com/questions/1050222/what-is-the-difference-between-concurrency-and-parallelism>
- Latency vs Throughput
 - <https://stackoverflow.com/questions/16718095/high-throughput-vs-low-latency-in-hdfs>